

Report: Data Wrangling

1. What kind of cleaning steps did you perform?

In order to collect the data I first needed to access my company's internal Redshift data warehouse. Using the python libraries sqlalchemy and psycopg2 (a postgresSQL driver) I queried five main tables representing Salesforce objects (Leads, Opportunities, Demos, Accounts, and Products). Demos are our main objects of interest with Leads, Accounts, Opportunities and Products providing further clarity into customer demographics and demo outcomes. Providing my user credentials and database information, I opened a connection to pass queries (which is later shutdown after the queries are completed".

An important concept to understand in analyzing sales data is the business process from Lead to Opportunity and how the different entities are represented. The stage of the sales cycle we are interested in examining will determine how the different tables are joined. The typical sales process is depicted below:

[Lead] → MQL → Sales Accepted Lead → Sales Qualified Lead → [Opportunity] → [Customer]

Along the way from Prospect to Customer, at least 3+ entities can be created when the individual enters the system as a Prospect, engages with our sales team, is converted to an Opportunity which is connected to an Account and the individual (now represented as a Contact). Each object will have a number of standard and custom editable fields that can be easily created to enrich a company's insight into the individual, deal, or company.

Given how frequently the metadata and schemas in Salesforce change in the start-up, I needed to query for all the relevant columns and fields and export csv samples for a visual inspection of the available data names, types and quality. Using the summaries, I manually constructed a data catalog showing the objects, related fields, the data warehouse names, the new names, necessary data transformations as well as possible data quality issues. After completing the first round of checks and evaluations and labeling fields which could be used for prediction or labeling, I created strings of candidate fields to subset the queried tables.

A	B	C	D	E	F	G	H	I	J	K	L
Original API Name	New Name	Type	select	new_names	Transformation	Transformation 2	Transformation 3	Use for model?	Use for EDA?	Data Issues	Type
email	email	Lead_PersonalInformation	email	email__Lead_PersonalInformation	Create new column - filled in non email/thermal/etc email domain		String	N	Y	Can be false "xyz", "123" Can be null	Lead: Personal Information
firstname	firstname	Lead_PersonalInformation	firstname	firstname__Lead_PersonalInformation	Create new column - filled in first name		String	N	Y	Can be false "xyz", "123" Can be null	Lead: Personal Information
lastname	lastname	Lead_PersonalInformation	lastname	lastname__Lead_PersonalInformation	Create new column - filled in last name		String	N	Y	Can be false "xyz", "123" Can be null	Lead: Personal Information
title	title	Lead_PersonalInformation	title	title__Lead_PersonalInformation	Create new column - filled in title		String	Y	Y	Not always filled	Lead: Personal Information
customer_type_c	customer_type	Lead_LeadCompanyInformation	customer_type_c	customer_type__Lead_LeadCompanyInformation	Create new column - filled in non-dummy		String	Y	Y	Can be null	Lead: Lead Company Information
company	company	Lead_LeadCompanyInformation	company	company__Lead_LeadCompanyInformation	Create new column - filled in company		String	N	Y	Not filled or false	Lead: Lead Company Information
street	street	Lead_LeadCompanyInformation	street	street__Lead_LeadCompanyInformation			Address	N	Y		Lead: Lead Company Information
city	city	Lead_LeadCompanyInformation	city	city__Lead_LeadCompanyInformation			Address	N	Y		Lead: Lead Company Information
state	state	Lead_LeadCompanyInformation	state	state__Lead_LeadCompanyInformation			Address	N	Y		Lead: Lead Company Information
country	country	Lead_LeadCompanyInformation	country	country__Lead_LeadCompanyInformation			Address	N	Y		Lead: Lead Company Information
linkedin_page_c	linkedinPage	Lead_MarketingInformation	linkedin_page_c	linkedinPage__Lead_MarketingInformation			String	Y	Y	Picklist values, can be null	Lead: Marketing Information
marketing_channel_c	marketingChannel	Lead_MarketingInformation	marketing_channel_c	marketingChannel__Lead_MarketingInformation			String	Y	Y	Can be null	Lead: Marketing Information
landing_page_c	landingPage	Lead_MarketingInformation	landing_page_c	landingPage__Lead_MarketingInformation	Create new column - Has a linkedin page		String	Y	Y	Can be null & not always filled up	Lead: Marketing Information
landing_page_url_c	landingPageUrl	Lead_MarketingInformation	landing_page_url_c	landingPageUrl__Lead_MarketingInformation			String	Y	Y	Can be null & not always filled up	Lead: Marketing Information
google_campaign_c	googleCampaign	Lead_MarketingInformation	google_campaign_c	googleCampaign__Lead_MarketingInformation			String	Y	Y	Optional	Lead: Marketing Information
leadsource	leadsource	Lead_MarketingInformation	leadsource	leadsource__Lead_MarketingInformation			String	Y	Y	Can have nulls	Lead: Marketing Information
converteddate	convertedDate	Lead_ConversionInformation	converteddate	convertedDate__Lead_ConversionInformation	Standardize date		Date	N	Y	3/19/2018	Lead: Conversion Information
status	statusReason	Lead_ConversionInformation	status	status__Lead_ConversionInformation	Filter out duplicates		Text	N	Y	Can have duplicates - but is their	Lead: Conversion Information
id	PK_LeadID	Lead_ConversionInformation	id	PK_LeadID__Lead_ConversionInformation			String	N	Y	Can have leads that are open	Lead: Conversion Information
convertedaccountid	PK_LeadtoAccount	Lead_ImportantJoinKey	convertedaccountid	PK_LeadtoAccount__Lead_ImportantJoinKey			String	N	Y		Lead: Important Join Key
convertedcontactid	PK_LeadtoContact	Lead_ImportantJoinKey	convertedcontactid	PK_LeadtoContact__Lead_ImportantJoinKey			String	N	Y		Lead: Important Join Key
convertedopportunityid	PK_LeadtoOpportunity	Lead_ImportantJoinKey	convertedopportunityid	PK_LeadtoOpportunity__Lead_ImportantJoinKey			String	N	Y		Lead: Important Join Key
ownerid	PK_LeadtoOwner	Lead_ImportantJoinKey	ownerid	PK_LeadtoOwner__Lead_ImportantJoinKey			String	N	Y		Lead: Important Join Key
createdbyid	createdbyid	Lead_ImportantSystemInfo	createdbyid	createdbyid__Lead_ImportantSystemInfo	Need to standardize to date		String	N	Y		Lead: Important System Info
createddate	createddate	Lead_ImportantSystemInfo	createddate	createddate__Lead_ImportantSystemInfo			Date	Y	Y		Lead: Important System Info
duplicate_lead_c	duplicateLead	Lead_ImportantSystemInfo	duplicate_lead_c	duplicateLead__Lead_ImportantSystemInfo	Need to filter out duplicate leads		Boolean	N	Y	Useless field	Lead: Important System Info
isconverted	isconverted	Lead_ImportantSystemInfo	isconverted	isconverted__Lead_ImportantSystemInfo			Boolean	N	Y		Lead: Important System Info
released	released	Lead_ImportantSystemInfo	released	released__Lead_ImportantSystemInfo	Filter out		Boolean	N	Y		Lead: Important System Info

Once the data frames were subset, the columns were renamed, indicating the attribute, originating object, and type of attribute.

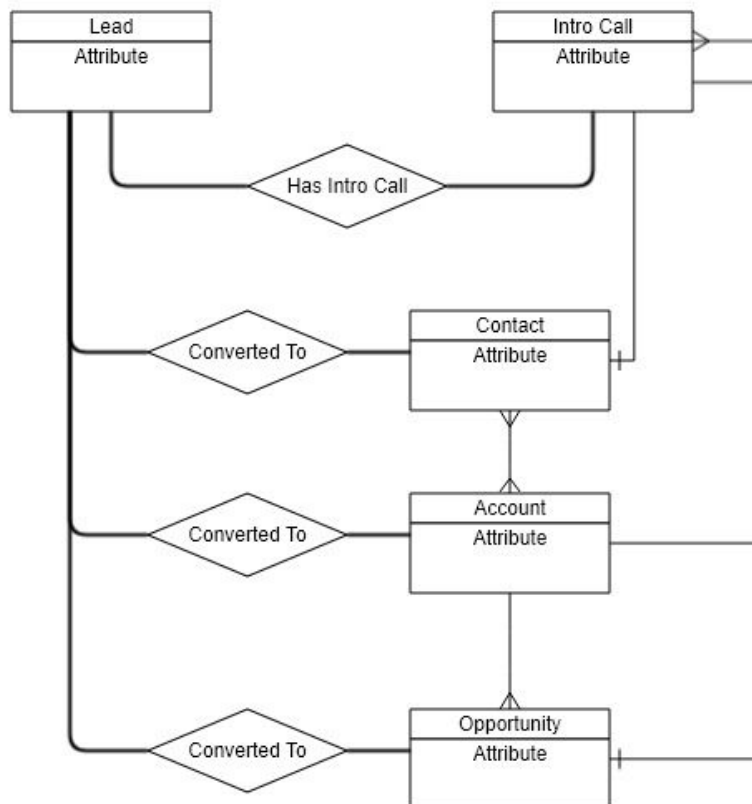
For example: "email__Lead_PersonalInformation" indicates the field describes Personal Information about a lead (in this case the email address) and came from the Lead dataframe (which was derived from the Lead table in the data warehouse). Another example: the field 'PK_OpptyID__Oppty_ImportantJoinKey' indicates that the field is the Opportunity ID, is meant to be used in a join, and is the primary key that describes an opportunity instance.

The naming convention and detailed data catalog is crucial for a few reasons: (1) similarly named fields could be duplicated across Salesforce objects without necessarily being related or exact - especially in the case of field mismatches; (2) real-time feedback was being given back to the data engineering team about data quality issues and solutions; (3) multiple data sets for exploration were being created and ideally we'd need clarity for the order of joins for 1:M and M:M relationships.


The subsetted data frames were then joined via the table/object keys identified during the data cataloging using merge. Two master data sets were created, masterDataSet where each demo call is a unique row (left joined by Leads, Opportunities, Accounts) and a masterDataSet_product where each product line item from the opportunity was left joined and enriched by the masterDataSet. The rationale for creating split data sets was to be able to accurately classify demo call outcomes but have the product data set

available in order to facilitate exploratory analysis around product purchases, SLA's, and support add-ons.

In sales analytics and sales operations, time stamping is crucial to: (1) estimating velocity of opportunity pipeline, (2) triaging individual opportunities for attention, (3)



meeting the agreed upon standard of performance. ensuring sales teams are



'2018-11-08T20:12:05.000Z', 'Q1-2015', '10/17/2014 17:09', '10/28/2014' are just four examples of the 10+ data columns that needed to be parsed and converted to a datetime object. I wrote a function that would take a dataframe, the target column to be parsed, name of a new column, and the date time pattern to pass in. The function `clean_dates` takes the specified columns, parses the timdate string, returns a datetime

object as the new column and deletes the old column.

The next import step was re-grouping the categorical features. After inspecting the different grouped columns and values, I remapped the values using dictionaries containing the new group values and deleted the old columns.

	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O	P	Q
1	Target__IntroCall_Outcome				rejectedReason__IntroCall_Outcome					status__Lead_ConversionInformation					trafficChannel__Lead_MarketingInformation		
2	Original Value	New Label			Original Value	New Label				Original Value	New Value				Original Value	New Value	
3	Attributed	Qualified			Company Too Small	Wrong_Demographic				Cancelled	Not Qualified				Affiliate	Affiliate	
4	Cancelled	Not Qualified			Does Not See Benefit of WalkMe	Not_Interested				Conference Rejuvenated	Open				Banner	Other	
5	No Show	Not Qualified			Existing Opportunity	Duplicate				Contacted	Open				Bing	Bing	
6	Qualified	Qualified			No Budget/Price Too High	Price_Too_High				Converted	Qualified				Biz Dev	Other	
7	Rejected	Rejected			No Commercial Influence	Not_Right_Person				Engaged	Open				BizLi	Other	
8	Rescheduling	Open			Not A Use Case Fit	Not_Interested				Finished Sequence	Open				Brand	Brand	
9	Scheduled	Open			Not Decision Maker	Not_Right_Person				Junk	Not Qualified				Bulk upload - R&D's bug	Other	
0					Other (please specify)	Other				Moved to SE	Open				Conference Emails	Email	
1					Project Fully Outsourced	Other				No Show	Not Qualified				Conferences	Event/Conference	
2					Startup - Too Expensive	Price_Too_High				No longer with company	Not Qualified				Conferences Lead Swaps	Event/Conference	
3					Too Few Users	Wrong_Demographic				Not Relevant	Not Qualified				Customer Engagement Event	Event/Conference	
4	"Attributed": "Qualified",				Wrong Source Version	Other				Nurture	Open				Email Nurturing	Email	
5	"Cancelled": "Not Qualified",				Wrong Timing	Not_Interested				Nurture (Outbound)	Open				Email Nurturing Conferences	Email	
6	"No Show": "Not Qualified",					Other				Open	Open				EmailMarketing	Email	
7	"Qualified": "Qualified",									Prospecting	Open				EmailNurturing	Email	

Fields with long text data were also dropped due to inexperience with NLP techniques.

Additional columns were created in order to measure the duration between various stages of leads, demos, and opportunities.

After creating all the necessary columns for exploratory analysis, additional id and demographic columns were deleted.

2. How did you deal with missing values, if any?

There are a couple categories of missing values:

	Value Missing? (Yes)(No)	
Value important? (Yes) (No)	Didn't include Ex:	Included
	Treat as additional enrichment	

Essentially, we would expect there to be missing values for some fields (given they wouldn't be populated until the user had progressed past a stage). For other fields the

columns had to be excluded because of poor practice and field enforcement within Salesforce.

For date columns I initially included a dummy date "1-1-1800" which required a fair amount of processing to produce reasonable charts and statistics on the avg cycle of a lead to an opportunity. On further advice I left the date fields blank. Given the data came from Salesforce, a majority of the date fields are created automatically by the CRM.

For demographic data like employee size I could impune the values based on the categorization of the account as accounts are stratified by employee size (i.e. an Account labeled Mid-Market vs Enterprise is supposed to fall within a certain employee size range, etc so even if the employee size is missing we can impune the value by choosing either the median of the range or the actual employee counts within the data).

3. Were there outliers, and how did you handle them?

The types of outliers that occurred and why were:

- Deal Age - Some opportunities were incredibly old because an opportunity had been created, dropped for some time, and then picked up by a different sales rep. Those outliers were kept given how that occur frequently due to sales rep turnover and should be reflected in efficiency metrics.